Assignment 5

Fall 2014
CS595 Web Science
Dr. Michael Nelson

Mathew Chaney

October 16, 2014

Contents

List of Figures

Listings

1 Question 1

1.1 Question

```
Explore the friendship paradox for your Twitter account.  Since
Twitter has directional links (i.e., "followers" and "following"),
we'll be investigating if the people you follow (Twitter calls these
people "friends") follow more people than you.  If you are following <
50 people, use my twitter account "phonedude_mln" instead of your own.

Create a graph of the number of friends (y-axis) and the friends sorted
by number of friends (x-axis).  (The friends don't need to be labeled on
the x-axis as "Bob", "Mary", etc. -- just 1, 2, 3 ...)  In other words,
if you have 100 friends your x-axis will be 1..101 (100 + you), and the
y-axis value will be number of friends that each of those friends has.
The friend with the lowest number of friends will be first and the friend
with the highest number of friends will be last.

Do include yourself in the graph and label yourself accordingly.  Compute
the mean, standard deviation, and median of the number of friends that
your friends have.

The appropriate part of the Twitter API to use is:

https://dev.twitter.com/rest/reference/get/friends/list
```

1.2 Answer

Using Dr. Michael Nelson's Twitter account and the Twitter API [1], specifically the GET friends/list [2] request, all of Dr. Nelson's Twitter friends were obtained and saved to the file called `friends`. This method also uses the API's paginating scheme: when there are a large number of results for a query, the API will send a cursor index to show that there are more results to process and that more requests are needed. The code to do this is in Listing 1.

```python
38  def get_friends(screen_name):
39      print("Getting {}'s friends".format(screen_name))
40      friends = []
41      next_cursor = -1
42      limit, reset = get_limit()
43      while True:
44          if limit == 0:
45              limit, reset = wait_for_reset(reset)
46          response = requests.get(SEARCH_URI, params={'screen_name':screen_name,
47                  'cursor':next_cursor, 'count':200}, auth=OAUTH)
48          if not response:
49              print("Bad response: {}".format(response.reason))
50              return []
51          limit = limit - 1
52          data = json.loads(response.text)
53          print("Got {} friends, limit: {}".format(len(data['users']), limit))
54          next_cursor = data['next_cursor']
55          friends.extend(data['users'])
56          if next_cursor == 0:
57              return [friend['screen_name'] for friend in friends]
```

Listing 1: Getting the friends list

To reduce the impact of high HTTP traffic, the Twitter API rate-limits most requests – the one needed to obtain a user's friends list has a limit of fifteen message per fifteen minutes. Any requests received from a user or service that has reached the limit will be denied. To ensure no HTTP requests are sent after the limit has been reached the script will sleep until the limit resets. This is accomplished using Python's `time` package [3] and the methods shown in Listing 2.

```python
24  def get_limit():
25      response = requests.get(LIMIT_URI, params={'resources':'friends'}, auth=OAUTH)
26      data = json.loads(response.text)
27      remaining = data['resources']['friends']['/friends/list']['remaining']
28      reset     = data['resources']['friends']['/friends/list']['reset']
29      return remaining, reset
30
31  def wait_for_reset(reset):
32      naptime = reset - time.time() + 5
33      print("Limit reached, sleeping for {}".format(naptime))
34      time.sleep(naptime)
35      print("Time to get up and go to work") # https://www.youtube.com/watch?v=ChP6SfmZVSE
36      return get_limit()
```

Listing 2: Waiting for API request limit reset

The `get_limit` method uses the API to find the number of available requests remaining for the GET friends/list method and also the time at which the limit will reset, received as seconds since the Unix epoch [4]. This method, combined with the `wait_for_reset` method, allowed the script to restart after an interruption and only require waiting for the appropriate amount of time. The sleep time was extended by 5 seconds to allow for a small buffer in case of mathematical errors.

The friends of Dr. Nelson's friends were then obtained with the same `get_friends` method from Listing 1 and stored in a file called `friend_counts`, each on a single line preceeded by their friend count. All of these operations were controlled by a main method, which is shown in Listing 3.

```python
59  if __name__ == '__main__':
60      # get list of MLN's friends from previously run get_friends method
61      with open('friends') as infile:
62          friends = [line.strip() for line in infile]
63
64      # initialize map of MLN's friends' friend counts
65      friend_counts = {friend:0 for friend in friends}
66
67      # get MLN's friends' friend counts which have already been processed
68      with open('friend_counts') as infile:
69          for line in infile:
70              count, friend = line.strip().split(' ')
71              friend_counts[friend] = count
72
73      # process the rest
74      for friend, count in friend_counts.iteritems():
75          if count > 0 and friend != 'phonedude_mln':
76              friends.remove(friend)
77      with open('friend_counts', 'a') as outfile:
78          for friend in friends:
79              friend_list = get_friends(friend)
80              if friend_list is None:
81                  continue
82              outfile.write('{} {}\n'.format(len(friend_list), friend))
```

Listing 3: Main method

4

The `friend_counts` file was ordered in place with the Unix command in Listing 4.

```
1  [mchaney@mchaney-l a5]$ cat friend_counts | sort -g -o friend_counts
```

Listing 4: Sort command

This file was then processed by the R script shown in Listing 5 to produce the graph in Figure 1

```
1  #! /usr/bin/Rscript
2
3  # read data
4  data <- read.table('friend_counts')
5  x <- seq(1, length(data$V1))
6  y <- data$V1
7
8  # get notable values
9  mln_idx <- grep("phonedude_mln", data$V2)
10 med_val <- median(data$V1)
11 med_idx <- which(abs(y - med_val) == min(abs(y - med_val)))
12 mean_val <- mean(data$V1)
13 mean_idx <- which(abs(y - mean_val) == min(abs(y - mean_val)))
14 std_dev <- sd(data$V1)
15
16 # draw the graph
17 pdf("friend_plot.pdf")
18 plot(x, y, type="l", log="y", pch=19, main="Dr. Nelson's Friends' Friends",
19     ylab="Number of Friends", xlab="Index of Friend")
20
21 # illustrate points of interest
22 abline(h=data$V1[mln_idx], col="red")
23 abline(h=data$V1[med_idx], col="blue")
24 abline(h=data$V1[mean_idx], col="darkolivegreen3")
25 abline(h=med_val + std_dev, col="purple")
26
27 # The Legend of the Data
28 legend(x=82, y=5, cex=0.8, lty=c(1, 1),
29     col=c("red", "blue", "darkolivegreen3", "white", "purple"),
30     c(paste("Nelson: ", data$V1[mln_idx]), paste("median: ", med_val),
31         paste("mean: ", format(round(mean_val, 4), nsmall = 4)),
32         paste("std dev: ", format(round(std_dev, 4), nsmall = 4)),
33         paste("median + 1 std dev: ", format(round(med_val + std_dev, 4), nsmall = 4))))
34 dev.off()
```

Listing 5: Graph Creation Script

The median, mean and standard deviation were all calculated, with the median, mean and median plus one standard deviation plotted as horizontal lines that intersect the data plot at their y-values. Only a single line was drawn for the standard deviation because the lower-end value was negative, and thus off the graph.


Figure 1: The Friendship Graph

2 References

[1] Twitter, Inc. Twitter API: Overview. https://dev.twitter.com/overview/api/, 2014.

[2] Twitter Inc. Twitter API: GET friends/list. https://dev.twitter.com/rest/reference/get/friends/list/, 2014.

[3] The Python Software Foundation. Python time module. https://docs.python.org/2/library/time.html, 2014.

[4] Community Wiki, Stack Overflow. Why is 1/1/1970 the "epoch time"? http://stackoverflow.com/questions/1090869/why-is-1-1-1970-the-epoch-time, Last edited June 23, 2011.