

Assignment 8

Fall 2014

CS595 Web Science

Dr. Michael Nelson

Mathew Chaney

December 2, 2014

Contents

1	Question 1	4
1.1	Question	4
1.2	Answer	4
2	Question 2	9
2.1	Question	9
2.2	Answer	9
3	Question 3	12
3.1	Question	12
3.2	Answer	12
4	Question 4	14
4.1	Question	14
4.2	Answer	14
5	Question 5	16
5.1	Question	16
5.2	Answer	16
6	Appendix A	20
7	References	28

Listings

1	referenced variables in get_uris.py	4
2	main for get_uris.py	4
3	get_atom function	5
4	add_uri function	5
5	looping over the URIs	5
6	processing each blog	6
7	creating the blog data matrix	6
8	loading the data	7
9	building the master wordlist	7
10	writing the data	7
11	building the histogram	7
12	creating the dendrograms	9
13	creating the dendrograms	9
14	hcluster function	9
15	printclust function	10
16	drawdendrogram function	10
17	kclustering main	12
18	kcluster function	12
19	output of kclustering algorithm	13
20	main for scaledown	14
21	scaledown function	14
22	draw2d function	15
23	use of tfidf function	16
24	writing the data	16

25	tf idf functions	16
26	get_uris.py	20
27	matrix.py	21
28	clusters.py	23
29	scaledown output	27

1 Question 1

1.1 Question

Create a blog-term matrix. Start by grabbing 100 blogs; include:

```
http://f-measure.blogspot.com/  
http://ws-dl.blogspot.com/
```

and grab 98 more as per the method shown in class.

Use the blog title as the identifier for each blog (and row of the matrix). Use the terms from every item/title (RSS) or entry/title (Atom) for the columns of the matrix. The values are the frequency of occurrence. Essentially you are replicating the format of the "blogdata.txt" file included with the PCI book code. Limit the number of terms to the most "popular" (i.e., frequent) 500 terms, this is *after* the criteria on p. 32 (slide 7) has been satisfied.

Create a histogram of how many pages each blog has (e.g., 30 blogs with just one page, 27 with two pages, 29 with 3 pages and so on).

1.2 Answer

To complete this assignment, a blog word count matrix was required. To start off, a list of blog URIs was obtained using the method described in class, implemented as the `get_uris.py` script, which can be found in Appendix A, Listing 26. Two default blogs, F-Measure and the Old Dominion Web Science and Digital Libraries blogs, were added as defaults to the initial URI list and then, using the seed URI provided (Listing 1), the remaining 98 URIs from random blogs within the blogger.com family were added.

```
7 default = 'http://www.blogger.com/next-blog?navBar=true&blogID=3471633091411211117'  
8 must_haves = ['http://f-measure.blogspot.com/', 'http://ws-dl.blogspot.com/']
```

Listing 1: referenced variables in `get_uris.py`

The `get_uris` main function in Listing 2 was the driver that called the `get_atom` function (shown in Listing 3) to extract the atom [1] URIs from each blog and add them to the set of URIs with the `add_uri` function, shown in Listing 4.

```
27 if __name__ == '__main__':  
28     uris = set()  
29     with open('blog_uris', 'a') as outfile:  
30         if len(sys.argv) > 1 and sys.argv[1] == 'new':  
31             for must_have in must_haves:  
32                 uri = get_atom(must_have)  
33                 add_uri(uri, uris, outfile)  
34         else:  
35             with open('blog_uris') as infile:  
36                 [uris.add(line.strip()) for line in infile]  
37         while len(uris) < 100:  
38             uri = get_atom(default)  
39             add_uri(uri, uris, outfile)
```

Listing 2: main for `get_uris.py`

```

10 def get_atom(uri):
11     try:
12         r = requests.get(uri)
13     except Exception, e:
14         return None
15     soup = BeautifulSoup(r.text)
16     links = soup.find_all('link', {'type': 'application/atom+xml'})
17     if links:
18         return str(links[0]['href'])
19     return None

```

Listing 3: get_atom function

```

21 def add_uri(uri, uris, outfile):
22     if uri and uri not in uris:
23         uris.add(uri)
24         outfile.write(uri + '\n')
25     print len(uris), uri

```

Listing 4: add_uri function

After the full list of 100 URIs was obtained, page counts for each blog were extracted and saved to a file called `pagecounts` using the `matrix.py` script. This script is a modified version of `generatefeedvectors.py` from the book *Programming Collective Intelligence* [2] and can be found in full in Appendix A, Listing 27.

The code responsible for downloading the blogs and counting the words in each is shown in Listing 5, which calls the `get_titles`, `get_words` and `get_next` functions found in Listing 6. This code loops over the list of URIs that was obtained with the `get_uris.py` script (Listing 26), parses each entry, and extracts all the words in each entry's title. These word counts are then saved as a python dictionary to the hard drive for later use.

```

95 if __name__ == '__main__':
96     with open('blog_uris') as infile:
97         uris = [line.strip() for line in infile if line.strip()]
98         if len(sys.argv) == 2 and sys.argv[1] == 'get':
99             with futures.ThreadPoolExecutor(max_workers=8) as executor:
100                 uri_futures = [executor.submit(get_titles, uri) for uri in uris]
101                 for future in futures.as_completed(uri_futures):
102                     uri, title, subtitle, pages, wc = future.result()
103                     with open('wcs/' + md5.new(uri).hexdigest() + '.w') as out:
104                         out.write(title + ': ' + subtitle + '\t' + str(pages) + '\t')
105                     json.dump(wc, out)

```

Listing 5: looping over the URIs

```

9 def get_next(d):
10     for item in d.feed.links:
11         if item['rel'] == u'next':
12             return item['href']
13     return None
14
15 def get_words(text):
16     txt = re.compile(r'<[>]+>').sub('', text)
17     words = re.compile(r'^A-Z^a-z+').split(txt)
18     return [word.lower() for word in words if word != '']
19
20 def get_titles(uri):
21     print('processing {}'.format(uri))
22     next = uri
23     wc = {}
24     pages = 0
25     while next is not None:
26         d = feedparser.parse(next)
27         for e in d.entries:
28             words = get_words(e.title.encode('utf-8'))
29             for word in words:
30                 wc.setdefault(word, 0)
31                 wc[word] += 1
32             pages += 1
33             next = get_next(d)
34             print('next {}'.format(next))
35     title = d.feed.title.encode('utf-8')
36     subtitle = d.feed.subtitle[:50].encode('utf-8')
37     print('finished: {}: {}'.format(title, subtitle))
38     return uri, title, subtitle, pages, wc

```

Listing 6: processing each blog

The parsed results were then read by the code in Listing 7. This code used the `load_data` and `build_wordlist` functions in Listing 8 and 9 to read each of the blog word counts and then created four collections to organize them all:

1. **apcount**: A dictionary containing the count for all words combined
2. **wordcounts**: A dictionary containing each blog's individual word count
3. **pagecounts**: A dictionary containing each blog's page count
4. **wordlist**: A list containing all of the words found in each blog

```

107     apcount, wordcounts, pagecounts = load_data(uris)
108     wordlist = build_wordlist(apcount, uris)
109     if len(sys.argv) == 2 and sys.argv[1] == 'pages':
110         with open('pagecounts', 'w') as outfile:
111             outfile.write('blog\tpages\n')
112             for blog, pagecount in pagecounts.iteritems():
113                 outfile.write("{}" + blog.replace("\n", "") + "\n" + '\t' + str(
114                     pagecount) + '\n')
115     elif len(sys.argv) == 2 and sys.argv[1] == 'wc':
116         write_data('blogdata1.txt', wordlist, wordcounts)

```

Listing 7: creating the blog data matrix

```

50 def load_data(uris):
51     apcount = {}
52     wordcounts = {}
53     pagecounts = {}
54     for uri in uris:
55         with open('wcs/' + md5.new(uri).hexdigest()) as infile:
56             try:
57                 lines = infile.read().split('\t')
58                 title = lines[0]
59                 pages = int(lines[1])
60                 wc = json.loads(lines[2])
61             except Exception, e:
62                 print('*** {} generated an exception: {}'.format(uri, e))
63                 continue
64             wordcounts[title] = wc
65             pagecounts[title] = pages
66             for word, count in wc.items():
67                 apcount.setdefault(word, 0)
68                 apcount[word] += count
69     return apcount, wordcounts, pagecounts

```

Listing 8: loading the data

```

71 def build_wordlist(apcount, uris):
72     wordlist = []
73     for w, bc in sorted(apcount.items(), key=lambda x: x[1], reverse=True):
74         frac = float(bc) / len(uris)
75         if frac > 0.1 and frac < 0.5:
76             wordlist.append(w)
77     return wordlist

```

Listing 9: building the master wordlist

The code in Listing 24 then created the matrix using the `write_data` function using the data structures that store the blog word counts.

```

79 def write_data(filename, wordlist, wordcounts, form=lambda wc, word, wordcounts: wc[word]):
80     with open(filename, 'w') as out:
81         out.write('Blog')
82         for word in wordlist[:500]:
83             out.write('\t%s' % word)
84         out.write('\n')
85         for blog, wc in wordcounts.items():
86             print blog
87             out.write(blog)
88             for word in wordlist[:500]:
89                 if word in wc:
90                     out.write('\t{}'.format(form(wc, word, wordcounts)))
91                 else:
92                     out.write('\t0')
93             out.write('\n')

```

Listing 10: writing the data

To build a histogram showing the blog page counts, the `pagecounts` file was parsed by the R script in Listing 11 and saved as a pdf, which is shown in Figure 1.

```

1 #! /usr/bin/Rscript
2
3 data <- read.table("pagecounts", sep="\t", header=TRUE, comment.char="")
4 counts <- table(data$pages)
5 pdf("hist.pdf")
6 barplot(counts, ylab="Number of Blogs", xlab="Page Count", main="Page Count per Blog")
7 dev.off()

```

Listing 11: building the histogram

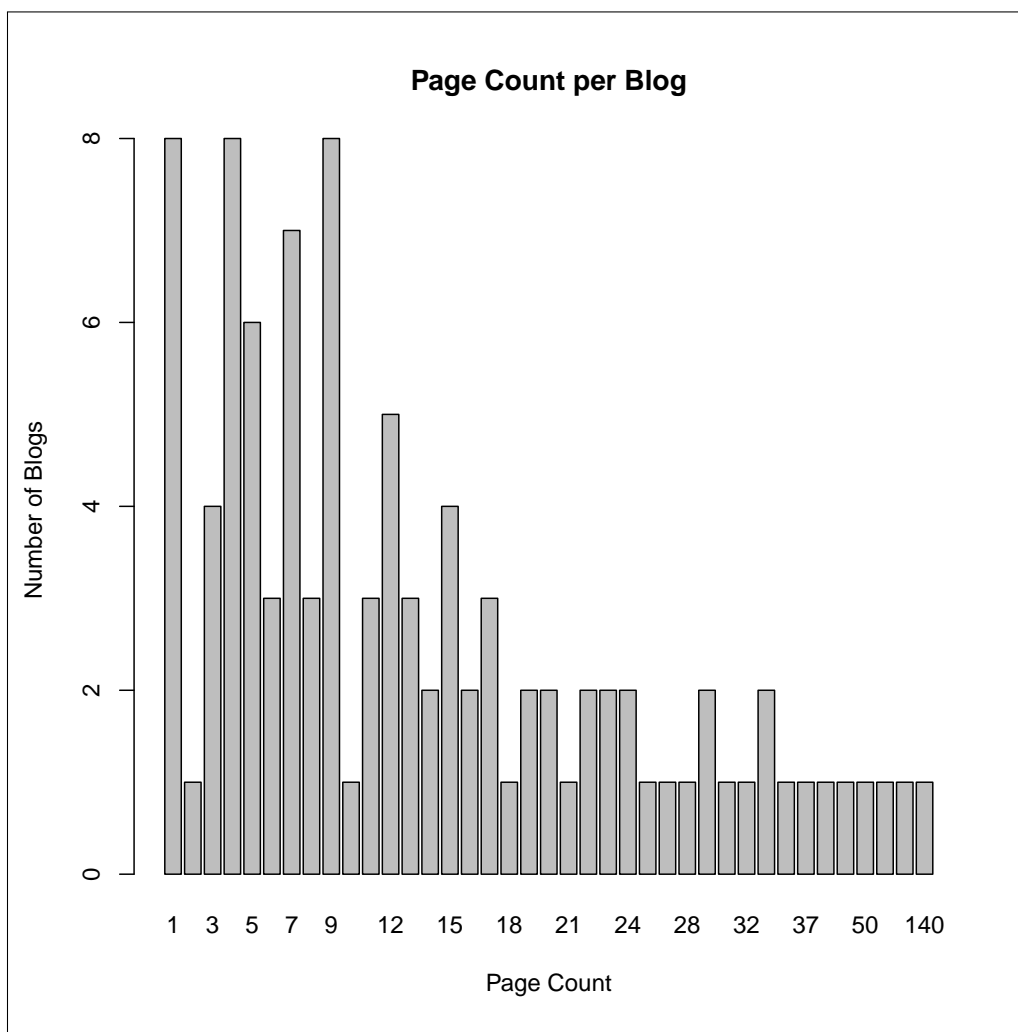


Figure 1: Page Count per Blog

2 Question 2

2.1 Question

Create an ASCII and JPEG dendrogram that clusters (i.e., HAC) the most similar blogs (see slides 12 & 13). Include the JPEG in your report and upload the ascii file to github (it will be too unwieldy for inclusion in the report).

2.2 Answer

The ascii and jpeg dendrograms were created using the code shown in Listing 12, which is modeled after the example from class.

```
286 blognames, words, data = readfile('q1/blogdata1.txt')
287 clust = hcluster(data)
288 with open('dendrogram.txt', 'w') as outfile:
289     stdout = sys.stdout
290     sys.stdout = outfile
291     printclust(clust, labels=blognames)
292     sys.stdout = stdout
293 drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
```

Listing 12: creating the dendrograms

The `readfile` function shown in Listing 13 was used to read the data that was compiled from Question 1 into memory where it is then processed by the `hcluster` function found in Listing 14 to produce the clustered representation of the blogs.

```
3 def readfile(filename):
4     lines=[line for line in file(filename)]
5
6     # First line is the column titles
7     colnames=lines[0].strip().split('\t')[1:]
8     rownames=[]
9     data=[]
10    for line in lines[1:]:
11        p=line.strip().split('\t')
12        # First column in each row is the rowname
13        rownames.append(p[0])
14        # The data for this row is the remainder of the row
15        data.append([float(x) for x in p[1:]])
16    return rownames, colnames, data
```

Listing 13: creating the dendrograms

```
48 def hcluster(rows, distance=pearson):
49     distances={}
50     currentclustid=-1
51
52     # Clusters are initially just the rows
53     clust=[bicluster(rows[i], id=i) for i in range(len(rows))]
54
55     while len(clust)>1:
56         lowestpair=(0,1)
57         closest=distance(clust[0].vec, clust[1].vec)
58
59         # loop through every pair looking for the smallest distance
60         for i in range(len(clust)):
61             for j in range(i+1, len(clust)):
62                 # distances is the cache of distance calculations
63                 if (clust[i].id, clust[j].id) not in distances:
64                     distances[(clust[i].id, clust[j].id)]=distance(clust[i].vec, clust[j].vec)
65
66                 d=distances[(clust[i].id, clust[j].id)]
67
68                 if d<closest:
69                     closest=d
```

```

70         lowestpair=(i,j)
71
72     # calculate the average of the two clusters
73     mergevec=[
74         (clust[lowestpair[0]].vec[i]+clust[lowestpair[1]].vec[i])/2.0
75         for i in range(len(clust[0].vec))
76     ]
77     # create the new cluster
78     newcluster=biclust(mergevec,left=clust[lowestpair[0]],
79                        right=clust[lowestpair[1]],
80                        distance=closest,id=currentclustid)
81
82     # cluster ids that weren't in the original set are negative
83     currentclustid-=1
84     del clust[lowestpair[1]]
85     del clust[lowestpair[0]]
86     clust.append(newcluster)
87
88     return clust[0]

```

Listing 14: hcluster function

The `printclust` function from Listing 15 prints the ascii dendrogram of the cluster object parameter to `sys.stdout`, which is redirected to write to a file with the code in Listing 12.

```

90 def printclust(clust, labels=None, n=0):
91     # indent to make a hierarchy layout
92     for i in range(n): print ' ',
93     if clust.id<0:
94         # negative id means that this is branch
95         print '- ',
96     else:
97         # positive id means that this is an endpoint
98         if labels==None: print clust.id
99         else: print labels[clust.id]
100
101     # now print the right and left branches
102     if clust.left!=None: printclust(clust.left, labels=labels, n=n+1)
103     if clust.right!=None: printclust(clust.right, labels=labels, n=n+1)

```

Listing 15: printclust function

The `drawdendrogram` function from Listing 16 creates a jpeg image of the cluster, which is shown in Figure 2.

```

122 def drawdendrogram(clust, labels, jpeg='clusters.jpg'):
123     # height and width
124     h=getheight(clust)*20
125     w=1200
126     depth=getdepth(clust)
127
128     # width is fixed, so scale distances accordingly
129     scaling=float(w-150)/depth
130
131     # Create a new image with a white background
132     img=Image.new('RGB',(w,h),(255,255,255))
133     draw=ImageDraw.Draw(img)
134
135     draw.line((0,h/2,10,h/2),fill=(255,0,0))
136
137     # Draw the first node
138     drawnode(draw,clust,10,(h/2),scaling,labels)
139     img.save(jpeg,'JPEG')

```

Listing 16: drawdendrogram function

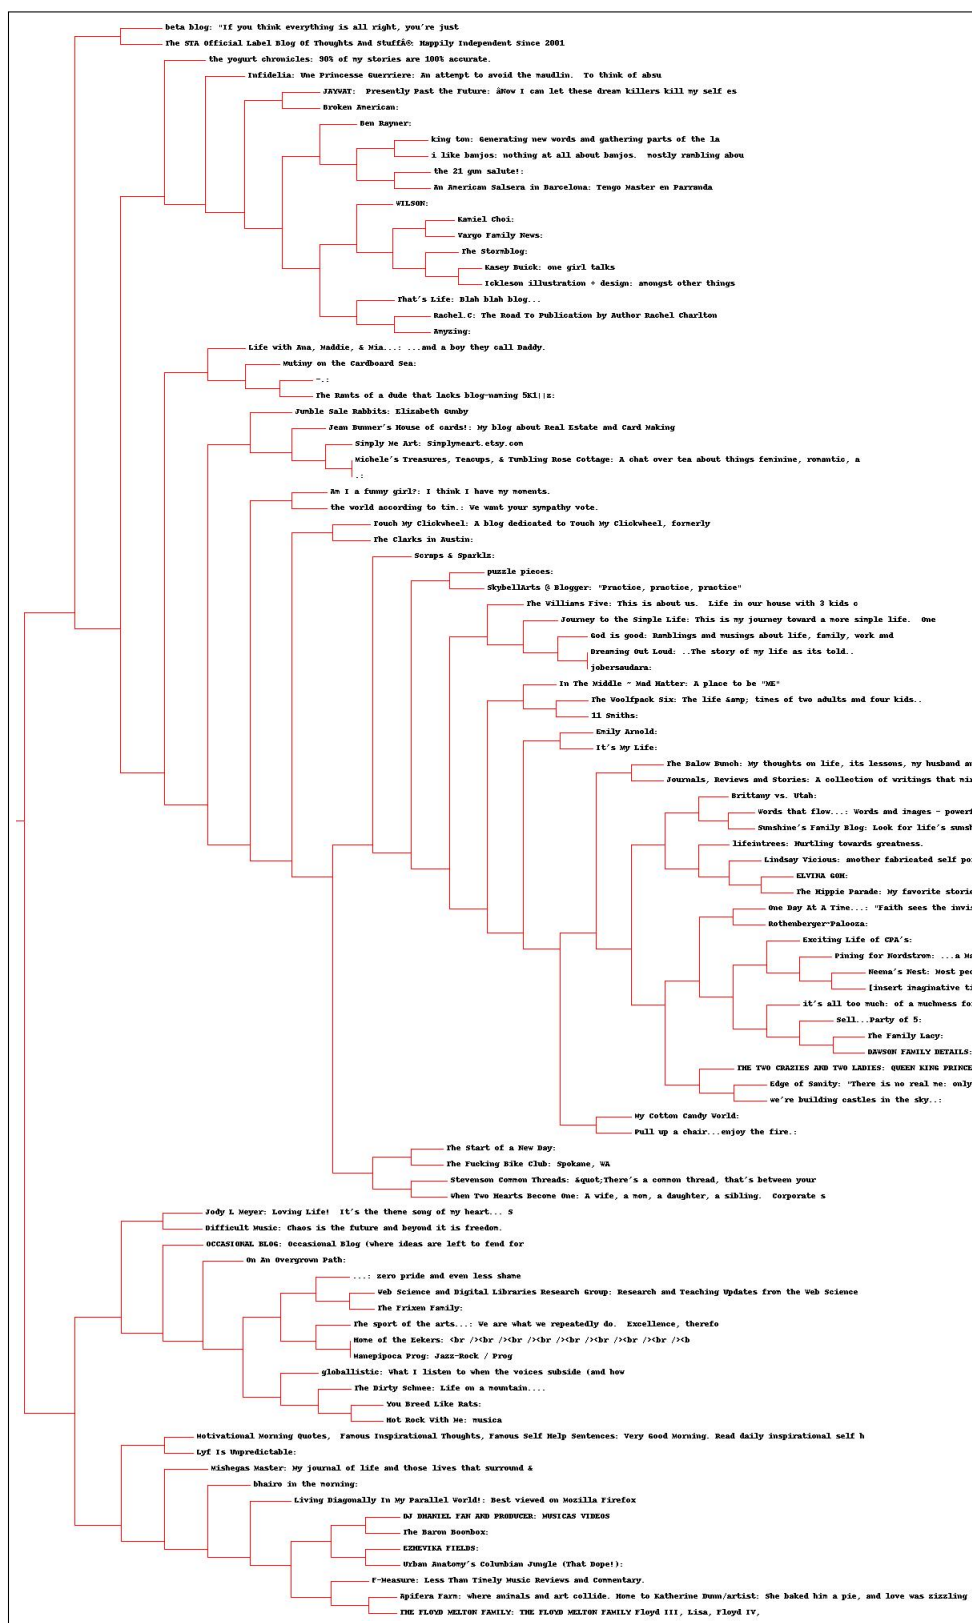


Figure 2: blog dendrogram

3 Question 3

3.1 Question

Cluster the blogs using K-Means, using $k=5, 10, 20$. (see slide 18).
How many iterations were required for each value of k ?

3.2 Answer

Using the code in Listing 17 kclustering was performed with values for $n = 5$, $n = 10$ and $n = 20$.
The main function calls the kcluster function, which is shown in Listing 18.

```
295 print "K=5"
296 kclust=kcluster(data, k=5)
297 print "K=10"
298 kclust=kcluster(data, k=10)
299 print "K=20"
300 kclust=kcluster(data, k=20)
```

Listing 17: kclustering main

```
174 def kcluster(rows, distance=pearson, k=4):
175     # Determine the minimum and maximum values for each point
176     ranges=[(min([row[i] for row in rows]), max([row[i] for row in rows]))
177             for i in range(len(rows[0]))]
178
179     # Create k randomly placed centroids
180     clusters=[[random.random()*(ranges[i][1]-ranges[i][0])+ranges[i][0]
181               for i in range(len(rows[0]))] for j in range(k)]
182
183     lastmatches=None
184     for t in range(100):
185         print 'Iteration %d' % t
186         bestmatches=[] for i in range(k)]
187
188         # Find which centroid is the closest for each row
189         for j in range(len(rows)):
190             row=rows[j]
191             bestmatch=0
192             for i in range(k):
193                 d=distance(clusters[i], row)
194                 if d<distance(clusters[bestmatch], row): bestmatch=i
195             bestmatches[bestmatch].append(j)
196
197         # If the results are the same as last time, this is complete
198         if bestmatches==lastmatches: break
199         lastmatches=bestmatches
200
201         # Move the centroids to the average of their members
202         for i in range(k):
203             avgs=[0.0]*len(rows[0])
204             if len(bestmatches[i])>0:
205                 for rowid in bestmatches[i]:
206                     for m in range(len(rows[rowid])):
207                         avgs[m]+=rows[rowid][m]
208                 for j in range(len(avgs)):
209                     avgs[j]/=len(bestmatches[i])
210                 clusters[i]=avgs
211
212     return bestmatches
```

Listing 18: kcluster function

The output is shown in Listing 19. As the output reads, a kcluster with $n = 5$ required nine iterations, $n = 10$ required four iterations and $n = 20$ also required four iterations.

```
1 Done with dendrograms
2 K=5
3 Iteration 0
4 Iteration 1
5 Iteration 2
6 Iteration 3
7 Iteration 4
8 Iteration 5
9 Iteration 6
10 Iteration 7
11 Iteration 8
12 K=10
13 Iteration 0
14 Iteration 1
15 Iteration 2
16 Iteration 3
17 K=20
18 Iteration 0
19 Iteration 1
20 Iteration 2
21 Iteration 3
```

Listing 19: output of kclustering algorithm

4 Question 4

4.1 Question

Use MDS to create a JPEG of the blogs similar to slide 29.
How many iterations were required?

4.2 Answer

With the code in Listing 20, multidimensional scaling (MDS) was used to create a two-dimensional visualization of the blog distance graph.

```
301 coords=scaledown(data)
302 draw2d(coords, blognames, jpeg='blogs2d.jpg')
```

Listing 20: main for scaledown

This main code calls the `scaledown` function, which is shown in Listing 21. The algorithm continues until the error factor stops decreasing, as shown in the output in Appendix A, Listing 29.

```
224 def scaledown(data,distance=pearson,rate=0.01):
225     n=len(data)
226
227     # The real distances between every pair of items
228     realdist=[[distance(data[i],data[j]) for j in range(n)]
229              for i in range(0,n)]
230
231     # Randomly initialize the starting points of the locations in 2D
232     loc=[[random.random(),random.random()] for i in range(n)]
233     fakedist=[[0.0 for j in range(n)] for i in range(n)]
234
235     lasterror=None
236     for m in range(0,1000):
237         # Find projected distances
238         for i in range(n):
239             for j in range(n):
240                 fakedist[i][j]=sqrt(sum([pow(loc[i][x]-loc[j][x],2)
241                                         for x in range(len(loc[i]))]))
242
243         # Move points
244         grad=[[0.0,0.0] for i in range(n)]
245
246         totalerror=0
247         for k in range(n):
248             for j in range(n):
249                 if j==k: continue
250                 # The error is percent difference between the distances
251                 if realdist[j][k] != 0:
252                     errorterm=(fakedist[j][k]-realdist[j][k])/realdist[j][k]
253
254                 # Each point needs to be moved away from or towards the other
255                 # point in proportion to how much error it has
256                 grad[k][0]+=((loc[k][0]-loc[j][0])/fakedist[j][k])*errorterm
257                 grad[k][1]+=((loc[k][1]-loc[j][1])/fakedist[j][k])*errorterm
258
259                 # Keep track of the total error
260                 totalerror+=abs(errorterm)
261         print totalerror
262
263         # If the answer got worse by moving the points, we are done
264         if lasterror and lasterror<totalerror: break
265         lasterror=totalerror
266
267         # Move each of the points by the learning rate times the gradient
268         for k in range(n):
269             loc[k][0]-=rate*grad[k][0]
270             loc[k][1]-=rate*grad[k][1]
271
272     return loc
```

Listing 21: scaledown function

The `scaledown` function returns the coordinates for each of the blogs in 2D space. This data was then used with the `draw2d` function in Listing 22, which produced the two-dimensional visualization created from the MDS algorithm, as shown in Figure 3.

```

274 def draw2d(data, labels, jpeg='mds2d.jpg'):
275     img=Image.new('RGB', (2000,2000), (255,255,255))
276     draw=ImageDraw.Draw(img)
277     for i in range(len(data)):
278         x=(data[i][0]+0.5)*1000
279         y=(data[i][1]+0.5)*1000
280         draw.text((x,y), labels[i], (0,0,0))
281     img.save(jpeg, 'JPEG')

```

Listing 22: draw2d function



Figure 3: MDS 2d visualization

5 Question 5

5.1 Question

Re-run question 2, but this time with proper TFIDF calculations instead of the hack discussed on slide 7 (p. 32). Use the same 500 words, but this time replace their frequency count with TFIDF scores as computed in assignment #3. Document the code, techniques, methods, etc. used to generate these TFIDF values. Upload the new data file to github.

Compare and contrast the resulting dendrogram with the dendrogram from question #2.

Note: ideally you would not reuse the same 500 terms and instead come up with TFIDF scores for all the terms and then choose the top 500 from that list, but I'm trying to limit the amount of work necessary.

5.2 Answer

To answer this question, `matrix.py` was modified to add the capability to calculate *Term Frequency Inverse Document Frequency (TF/IDF)*. The added functions for computing TF/IDF are found in Listing 25. These functions use the master word count dictionary (`wordcounts`) and each blog's individual word count (`wc`) for each of the words in the `wordlist` from Question 1/2 to compute the TF/IDF value for each word in the word list.

```
116     elif len(sys.argv) == 2 and sys.argv[1] == 'tfidf':
117         write_data('blogdata2.txt', wordlist, wordcounts, form=lambda wc, word,
                    wordcounts: tf(wc, word) * idf(wordcounts, word))
```

Listing 23: use of tfidf function

```
79 def write_data(filename, wordlist, wordcounts, form=lambda wc, word, wordcounts: wc[word]):
80     with open(filename, 'w') as out:
81         out.write('Blog')
82         for word in wordlist[:500]:
83             out.write('\t%s' % word)
84         out.write('\n')
85         for blog, wc in wordcounts.items():
86             print blog
87             out.write(blog)
88             for word in wordlist[:500]:
89                 if word in wc:
90                     out.write('\t{}'.format(form(wc, word, wordcounts)))
91                 else:
92                     out.write('\t0')
93             out.write('\n')
```

Listing 24: writing the data

```
40 def tf(wc, word):
41     return float(wc[word]) / float(sum(wc.values()))
42
43 def idf(wordcounts, word):
44     present = 0
45     for wc in wordcounts.values():
46         if word in wc:
47             present += 1
48     return math.log(len(wordcounts) / present, 2)
```

Listing 25: tf idf functions

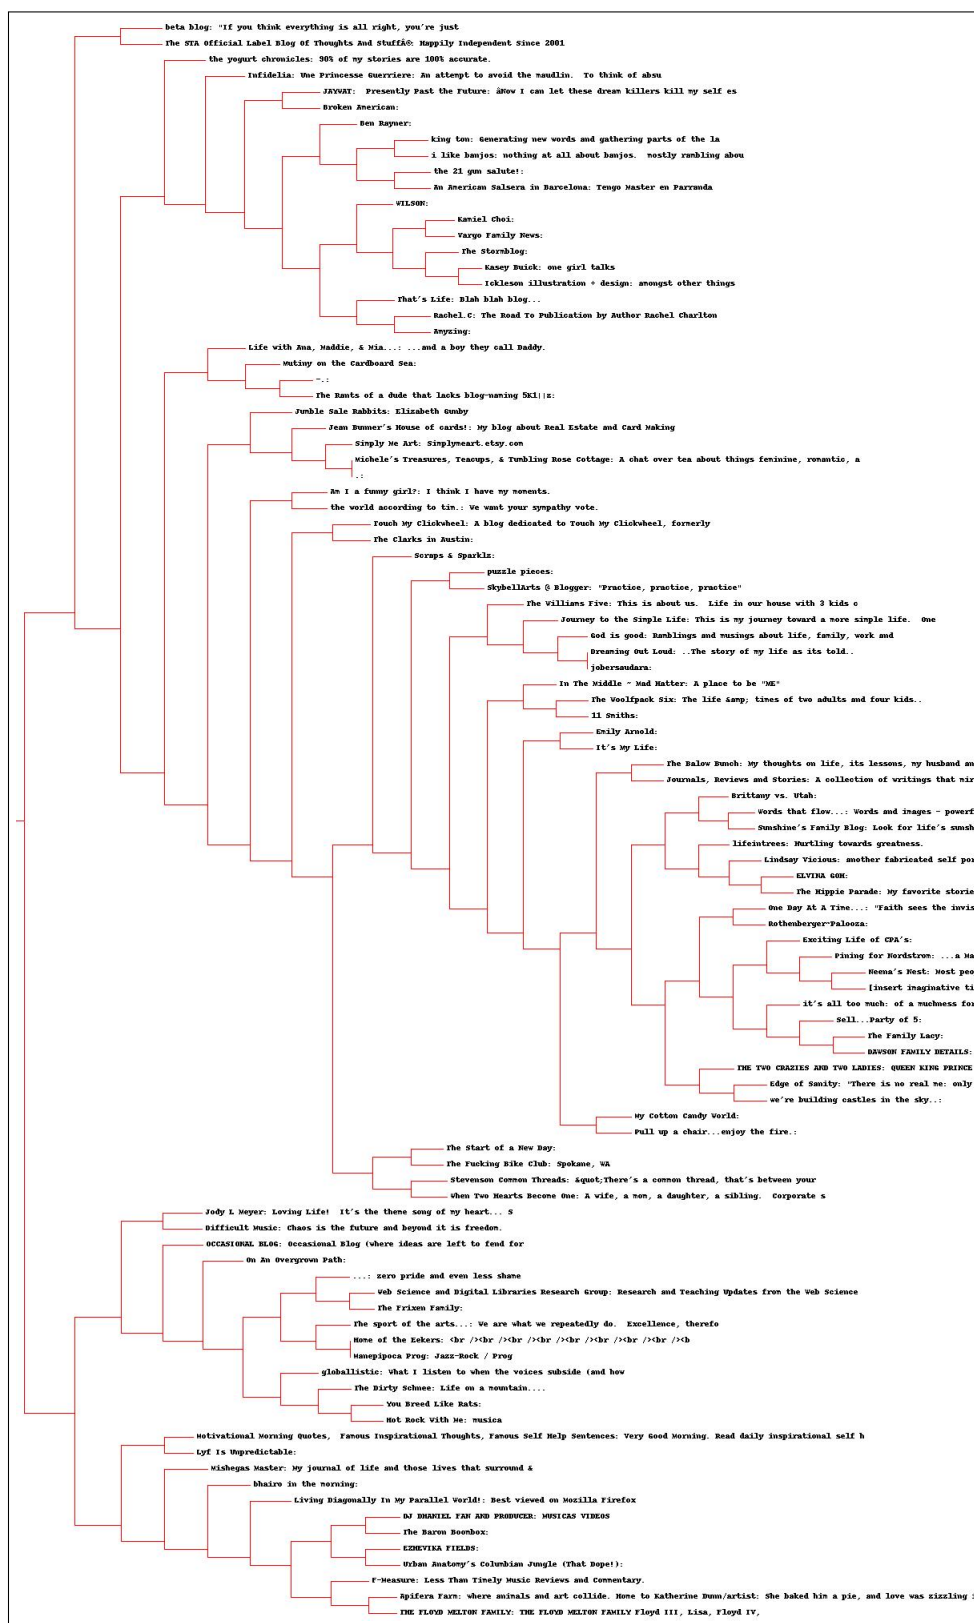
The same clustering was applied to the TF/IDF result matrix as was done in Question 2 and both images are displayed in Figures 4 and 5.

There are a many pairs that were found to be similar in both clusterings. For example, in both dendrograms, the *Web Science and Digital Libraries Research Group* blog is most similar to the *...: zero pride and even less shame* blog, the *DJ DHANIEL FAN AND PRODUCER* and *The Baron Boombox*, among a few others. In spite of this, the larger groupings do not appear very similar between the two clustings.

When examining the clusters on a larger scale, it seems that the TF/IDF dendrogram clustered blogs are subjectively more alike better than the raw count version did. Looking toward the bottom of the TF/IDF clustering image, one will notice a grouping of blogs that seem closely related to music: *F-Measure: Less Than timely Music Reviews and Commentary.*, *Urban Anatomy's Columbian Jungle (That Dope!)* which seems to be a melding of fashion and music commentary, *Ezhevika Fields* a blog where info and preview samples of "lost album samples of the past" can be found, whereas these blogs are not all grouped together in the raw count version. It seems that using TF/IDF as a basis for document comparison gives much more context than a raw term count (which doesn't take into consideration the rest of the document in which each word resides).

There is another cluster in the TF/IDF driven image that seems to contain family related blogs, with blogs like *Am I a funny girl?: I think I have my moments*, *The Frixen Family*, *When Two Hearts Become One: A wife, a mom, a daughter, a sibling.* and *The Clarks in Austin* all being related to a particular family and their everyday lives. Some of these blogs are close to each other in the raw count version, but they are not separated into their own distinct groups.

When looking at the overall structure of each dendrogram, it becomes apparent that there are more individual clusters grouped together in the TF/IDF version than are present in the raw count image, where there seems to be few small clusters and one mega-cluster in the center. This suggests that the TF/IDF algorithm is better at defining discrete subgroups within a larger context than a simple raw count driven dendrogram will produce.



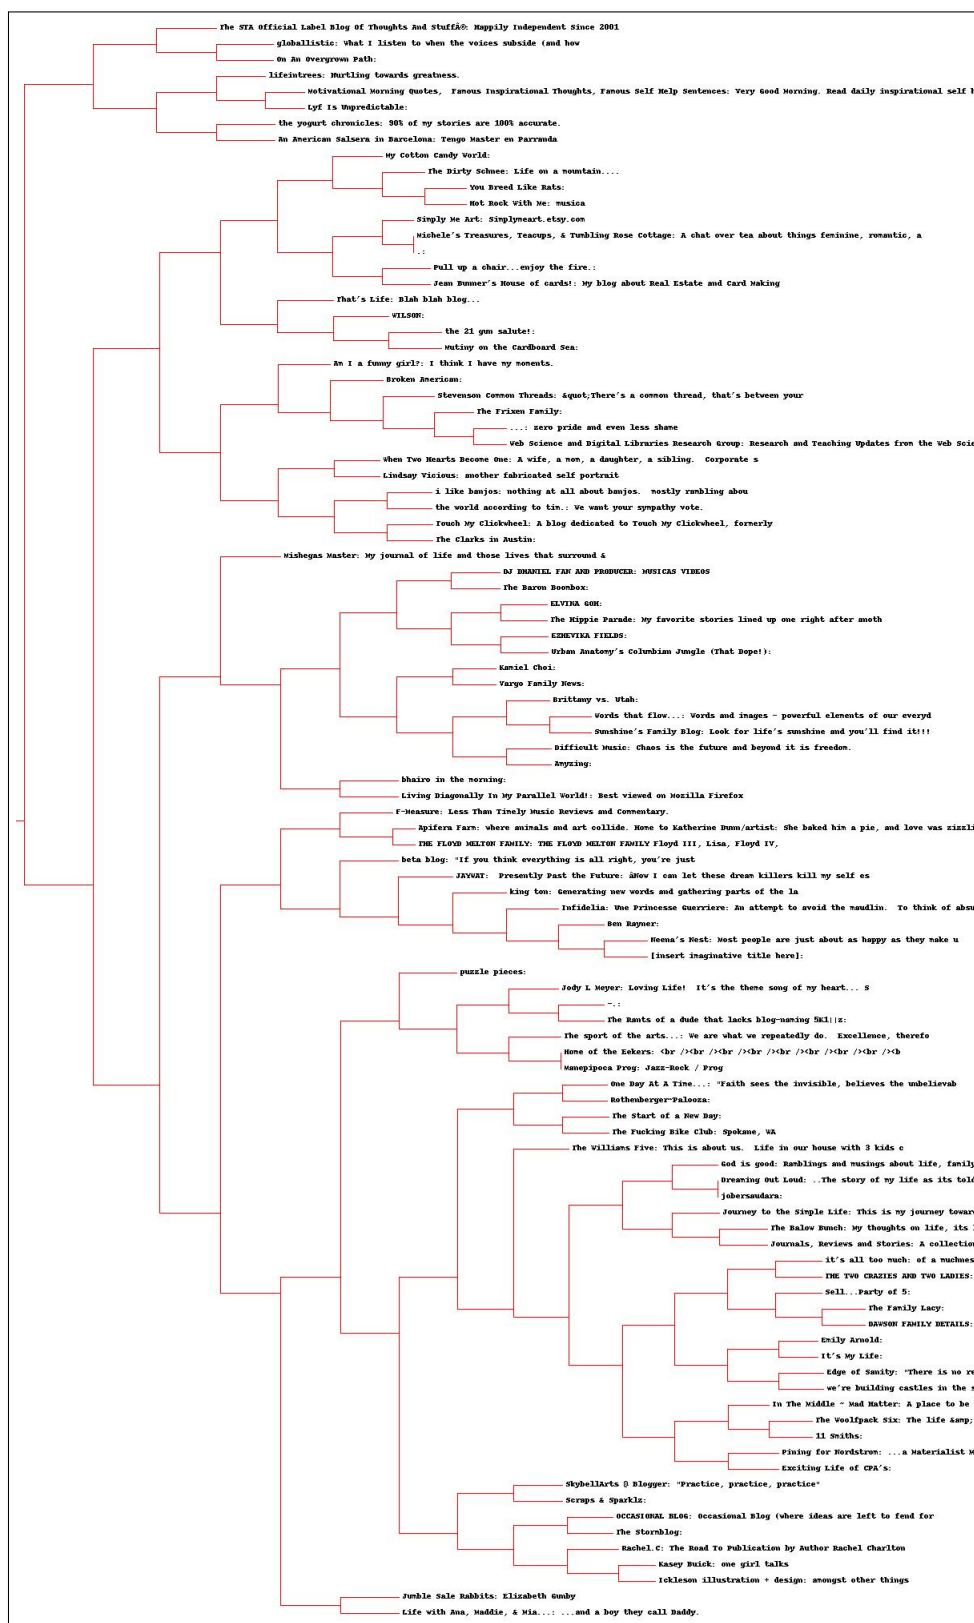


Figure 5: TF/IDF driven dendrogram

6 Appendix A

```
1 #!/usr/bin/env python
2
3 import requests
4 import sys
5 from bs4 import BeautifulSoup
6
7 default = 'http://www.blogger.com/next-blog?navBar=true&blogID=3471633091411211117'
8 must_haves = ['http://f-measure.blogspot.com/', 'http://ws-dl.blogspot.com/']
9
10 def get_atom(uri):
11     try:
12         r = requests.get(uri)
13     except Exception, e:
14         return None
15     soup = BeautifulSoup(r.text)
16     links = soup.find_all('link', {'type': 'application/atom+xml'})
17     if links:
18         return str(links[0]['href'])
19     return None
20
21 def add_uri(uri, uris, outfile):
22     if uri and uri not in uris:
23         uris.add(uri)
24         outfile.write(uri + '\n')
25         print len(uris), uri
26
27 if __name__ == '__main__':
28     uris = set()
29     with open('blog_uris', 'a') as outfile:
30         if len(sys.argv) > 1 and sys.argv[1] == 'new':
31             for must_have in must_haves:
32                 uri = get_atom(must_have)
33                 add_uri(uri, uris, outfile)
34         else:
35             with open('blog_uris') as infile:
36                 [uris.add(line.strip()) for line in infile]
37     while len(uris) < 100:
38         uri = get_atom(default)
39         add_uri(uri, uris, outfile)
```

Listing 26: get_uris.py

```

1 import feedparser
2 import futures
3 import math
4 import md5
5 import re
6 import sys
7 import json
8
9 def get_next(d):
10     for item in d.feed.links:
11         if item['rel'] == u'next':
12             return item['href']
13     return None
14
15 def get_words(text):
16     txt = re.compile(r'<[^>]+>').sub('', text)
17     words = re.compile(r'^A-Z^a-z+').split(txt)
18     return [word.lower() for word in words if word != '']
19
20 def get_titles(uri):
21     print('processing {}'.format(uri))
22     next = uri
23     wc = {}
24     pages = 0
25     while next is not None:
26         d = feedparser.parse(next)
27         for e in d.entries:
28             words = get_words(e.title.encode('utf-8'))
29             for word in words:
30                 wc.setdefault(word, 0)
31                 wc[word] += 1
32             pages += 1
33             next = get_next(d)
34             print('next {}'.format(next))
35         title = d.feed.title.encode('utf-8')
36         subtitle = d.feed.subtitle[:50].encode('utf-8')
37         print('finished: {}: {}'.format(title, subtitle))
38         return uri, title, subtitle, pages, wc
39
40 def tf(wc, word):
41     return float(wc[word]) / float(sum(wc.values()))
42
43 def idf(wordcounts, word):
44     present = 0
45     for wc in wordcounts.values():
46         if word in wc:
47             present += 1
48     return math.log(len(wordcounts) / present, 2)
49
50 def load_data(uris):
51     apcount = {}
52     wordcounts = {}
53     pagecounts = {}
54     for uri in uris:
55         with open('wcs/' + md5.new(uri).hexdigest()) as infile:
56             try:
57                 lines = infile.read().split('\n')
58                 title = lines[0]
59                 pages = int(lines[1])
60                 wc = json.loads(lines[2])
61             except Exception, e:
62                 print('*** {} generated an exception: {}'.format(uri, e))
63                 continue
64             wordcounts[title] = wc
65             pagecounts[title] = pages
66             for word, count in wc.items():
67                 apcount.setdefault(word, 0)
68                 apcount[word] += count
69     return apcount, wordcounts, pagecounts
70
71 def build_wordlist(apcount, uris):
72     wordlist = []
73     for w, bc in sorted(apcount.items(), key=lambda x: x[1], reverse=True):
74         frac = float(bc) / len(uris)
75         if frac > 0.1 and frac < 0.5:
76             wordlist.append(w)

```

```

77     return wordlist
78
79 def write_data(filename, wordlist, wordcounts, form=lambda wc, word, wordcounts: wc[word]):
80     with open(filename, 'w') as out:
81         out.write('Blog')
82         for word in wordlist[:500]:
83             out.write('\t%s' % word)
84         out.write('\n')
85         for blog, wc in wordcounts.items():
86             print blog
87             out.write(blog)
88             for word in wordlist[:500]:
89                 if word in wc:
90                     out.write('\t{}'.format(form(wc, word, wordcounts)))
91                 else:
92                     out.write('\t0')
93             out.write('\n')
94
95 if __name__ == '__main__':
96     with open('blog_uris') as infile:
97         uris = [line.strip() for line in infile if line.strip()]
98         if len(sys.argv) == 2 and sys.argv[1] == 'get':
99             with futures.ThreadPoolExecutor(max_workers=8) as executor:
100                 uri_futures = [executor.submit(get_titles, uri) for uri in uris]
101                 for future in futures.as_completed(uri_futures):
102                     uri, title, subtitle, pages, wc = future.result()
103                     with open('wcs/' + md5.new(uri).hexdigest(), 'w') as out:
104                         out.write(title + ': ' + subtitle + '\t' + str(pages) + '\t')
105                     json.dump(wc, out)
106             else:
107                 apcount, wordcounts, pagecounts = load_data(uris)
108                 wordlist = build_wordlist(apcount, uris)
109                 if len(sys.argv) == 2 and sys.argv[1] == 'pages':
110                     with open('pagecounts', 'w') as outfile:
111                         outfile.write('blog\tpages\n')
112                         for blog, pagecount in pagecounts.iteritems():
113                             outfile.write("'" + blog.replace("\n", "") + "'" + "\n" + '\t' + str(
                                pagecount) + '\n')
114                 elif len(sys.argv) == 2 and sys.argv[1] == 'wc':
115                     write_data('blogdata1.txt', wordlist, wordcounts)
116                 elif len(sys.argv) == 2 and sys.argv[1] == 'tfidf':
117                     write_data('blogdata2.txt', wordlist, wordcounts, form=lambda wc, word,
                                wordcounts: tf(wc, word) * idf(wordcounts, word))

```

Listing 27: matrix.py

```

1 from PIL import Image, ImageDraw
2
3 def readfile(filename):
4     lines=[line for line in file(filename)]
5
6     # First line is the column titles
7     colnames=lines[0].strip().split('\t')[1:]
8     rownames=[]
9     data=[]
10    for line in lines[1:]:
11        p=line.strip().split('\t')
12        # First column in each row is the rowname
13        rownames.append(p[0])
14        # The data for this row is the remainder of the row
15        data.append([float(x) for x in p[1:]])
16    return rownames, colnames, data
17
18
19 from math import sqrt
20
21 def pearson(v1,v2):
22     # Simple sums
23     sum1=sum(v1)
24     sum2=sum(v2)
25
26     # Sums of the squares
27     sum1Sq=sum([pow(v,2) for v in v1])
28     sum2Sq=sum([pow(v,2) for v in v2])
29
30     # Sum of the products
31     pSum=sum([v1[i]*v2[i] for i in range(len(v1))])
32
33     # Calculate r (Pearson score)
34     num=pSum-(sum1*sum2/len(v1))
35     den=sqrt((sum1Sq-pow(sum1,2)/len(v1))*(sum2Sq-pow(sum2,2)/len(v1)))
36     if den==0: return 0
37
38     return 1.0-num/den
39
40 class bicluster:
41     def __init__(self, vec, left=None, right=None, distance=0.0, id=None):
42         self.left=left
43         self.right=right
44         self.vec=vec
45         self.id=id
46         self.distance=distance
47
48 def hcluster(rows, distance=pearson):
49     distances={}
50     currentclustid=-1
51
52     # Clusters are initially just the rows
53     clust=[bicluster(rows[i], id=i) for i in range(len(rows))]
54
55     while len(clust)>1:
56         lowestpair=(0,1)
57         closest=distance(clust[0].vec, clust[1].vec)
58
59         # loop through every pair looking for the smallest distance
60         for i in range(len(clust)):
61             for j in range(i+1, len(clust)):
62                 # distances is the cache of distance calculations
63                 if (clust[i].id, clust[j].id) not in distances:
64                     distances[(clust[i].id, clust[j].id)]=distance(clust[i].vec, clust[j].vec)
65
66                 d=distances[(clust[i].id, clust[j].id)]
67
68                 if d<closest:
69                     closest=d
70                     lowestpair=(i,j)
71
72         # calculate the average of the two clusters
73         mergevec=[
74             (clust[lowestpair[0]].vec[i]+clust[lowestpair[1]].vec[i])/2.0
75             for i in range(len(clust[0].vec))]
76

```

```

77 # create the new cluster
78 newcluster=biclusterc(mergevec, left=clust[lowestpair[0]],
79                        right=clust[lowestpair[1]],
80                        distance=closest, id=currentclustid)
81
82 # cluster ids that weren't in the original set are negative
83 currentclustid -=1
84 del clust[lowestpair[1]]
85 del clust[lowestpair[0]]
86 clust.append(newcluster)
87
88 return clust[0]
89
90 def printclust(clust, labels=None, n=0):
91     # indent to make a hierarchy layout
92     for i in range(n): print ' ',
93     if clust.id < 0:
94         # negative id means that this is branch
95         print '- ',
96     else:
97         # positive id means that this is an endpoint
98         if labels==None: print clust.id
99         else: print labels[clust.id]
100
101 # now print the right and left branches
102 if clust.left!=None: printclust(clust.left, labels=labels, n=n+1)
103 if clust.right!=None: printclust(clust.right, labels=labels, n=n+1)
104
105 def getheight(clust):
106     # Is this an endpoint? Then the height is just 1
107     if clust.left==None and clust.right==None: return 1
108
109     # Otherwise the height is the same of the heights of
110     # each branch
111     return getheight(clust.left)+getheight(clust.right)
112
113 def getdepth(clust):
114     # The distance of an endpoint is 0.0
115     if clust.left==None and clust.right==None: return 0
116
117     # The distance of a branch is the greater of its two sides
118     # plus its own distance
119     return max(getdepth(clust.left), getdepth(clust.right))+clust.distance
120
121
122 def drawdendrogram(clust, labels, jpeg='clusters.jpg'):
123     # height and width
124     h=getheight(clust)*20
125     w=1200
126     depth=getdepth(clust)
127
128     # width is fixed, so scale distances accordingly
129     scaling=float(w-150)/depth
130
131     # Create a new image with a white background
132     img=Image.new('RGB', (w,h), (255,255,255))
133     draw=ImageDraw.Draw(img)
134
135     draw.line((0, h/2, 10, h/2), fill=(255,0,0))
136
137     # Draw the first node
138     drawnode(draw, clust, 10, (h/2), scaling, labels)
139     img.save(jpeg, 'JPEG')
140
141 def drawnode(draw, clust, x, y, scaling, labels):
142     if clust.id < 0:
143         h1=getheight(clust.left)*20
144         h2=getheight(clust.right)*20
145         top=y-(h1+h2)/2
146         bottom=y+(h1+h2)/2
147         # Line length
148         ll=clust.distance*scaling
149         # Vertical line from this cluster to children
150         draw.line((x, top+h1/2, x, bottom-h2/2), fill=(255,0,0))
151
152         # Horizontal line to left item
153         draw.line((x, top+h1/2, x+ll, top+h1/2), fill=(255,0,0))

```



```

154
155     # Horizontal line to right item
156     draw.line((x,bottom-h2/2,x+11,bottom-h2/2),fill=(255,0,0))
157
158     # Call the function to draw the left and right nodes
159     drawnode(draw,clust.left,x+11,top+h1/2,scaling,labels)
160     drawnode(draw,clust.right,x+11,bottom-h2/2,scaling,labels)
161 else:
162     # If this is an endpoint, draw the item label
163     draw.text((x+5,y-7),labels[clust.id],(0,0,0))
164
165 def rotatematrix(data):
166     newdata=[]
167     for i in range(len(data[0])):
168         newrow=[data[j][i] for j in range(len(data))]
169         newdata.append(newrow)
170     return newdata
171
172 import random
173
174 def kcluster(rows,distance=pearson,k=4):
175     # Determine the minimum and maximum values for each point
176     ranges=[(min([row[i] for row in rows]),max([row[i] for row in rows]))
177             for i in range(len(rows[0]))]
178
179     # Create k randomly placed centroids
180     clusters=[[random.random()*(ranges[i][1]-ranges[i][0])+ranges[i][0]
181               for i in range(len(rows[0]))] for j in range(k)]
182
183     lastmatches=None
184     for t in range(100):
185         print 'Iteration %d' % t
186         bestmatches=[] for i in range(k)]
187
188         # Find which centroid is the closest for each row
189         for j in range(len(rows)):
190             row=rows[j]
191             bestmatch=0
192             for i in range(k):
193                 d=distance(clusters[i],row)
194                 if d<distance(clusters[bestmatch],row): bestmatch=i
195             bestmatches[bestmatch].append(j)
196
197         # If the results are the same as last time, this is complete
198         if bestmatches==lastmatches: break
199         lastmatches=bestmatches
200
201         # Move the centroids to the average of their members
202         for i in range(k):
203             avgs=[0.0]*len(rows[0])
204             if len(bestmatches[i])>0:
205                 for rowid in bestmatches[i]:
206                     for m in range(len(rows[rowid])):
207                         avgs[m]+=rows[rowid][m]
208                 for j in range(len(avgs)):
209                     avgs[j]/=len(bestmatches[i])
210                 clusters[i]=avgs
211
212     return bestmatches
213
214 def tanamoto(v1,v2):
215     c1,c2,shr=0,0,0
216
217     for i in range(len(v1)):
218         if v1[i]!=0: c1+=1 # in v1
219         if v2[i]!=0: c2+=1 # in v2
220         if v1[i]!=0 and v2[i]!=0: shr+=1 # in both
221
222     return 1.0-(float(shr)/(c1+c2-shr))
223
224 def scaledown(data,distance=pearson,rate=0.01):
225     n=len(data)
226
227     # The real distances between every pair of items
228     realdist=[[distance(data[i],data[j]) for j in range(n)]
229               for i in range(0,n)]
230

```

```

231 # Randomly initialize the starting points of the locations in 2D
232 loc=[[random.random(),random.random()] for i in range(n)]
233 fakedist=[[0.0 for j in range(n)] for i in range(n)]
234
235 lasterror=None
236 for m in range(0,1000):
237     # Find projected distances
238     for i in range(n):
239         for j in range(n):
240             fakedist[i][j]=sqrt(sum([pow(loc[i][x]-loc[j][x],2)
241                                     for x in range(len(loc[i]))]))
242
243     # Move points
244     grad=[[0.0,0.0] for i in range(n)]
245
246     totalerror=0
247     for k in range(n):
248         for j in range(n):
249             if j==k: continue
250             # The error is percent difference between the distances
251             if realdist[j][k] != 0:
252                 errorterm=(fakedist[j][k]-realdist[j][k])/realdist[j][k]
253
254             # Each point needs to be moved away from or towards the other
255             # point in proportion to how much error it has
256             grad[k][0]+=((loc[k][0]-loc[j][0])/fakedist[j][k])*errorterm
257             grad[k][1]+=((loc[k][1]-loc[j][1])/fakedist[j][k])*errorterm
258
259             # Keep track of the total error
260             totalerror+=abs(errorterm)
261     print totalerror
262
263     # If the answer got worse by moving the points, we are done
264     if lasterror and lasterror<totalerror: break
265     lasterror=totalerror
266
267     # Move each of the points by the learning rate times the gradient
268     for k in range(n):
269         loc[k][0]-=rate*grad[k][0]
270         loc[k][1]-=rate*grad[k][1]
271
272     return loc
273
274 def draw2d(data,labels, jpeg='mds2d.jpg'):
275     img=Image.new('RGB',(2000,2000),(255,255,255))
276     draw=ImageDraw.Draw(img)
277     for i in range(len(data)):
278         x=(data[i][0]+0.5)*1000
279         y=(data[i][1]+0.5)*1000
280         draw.text((x,y),labels[i],(0,0,0))
281     img.save(jpeg,'JPEG')
282
283 import sys
284
285 if __name__ == '__main__':
286     blognames, words, data = readfile('q1/blogdata1.txt')
287     clust = hcluster(data)
288     with open('dendrogram.txt', 'w') as outfile:
289         stdout = sys.stdout
290         sys.stdout = outfile
291         printclust(clust, labels=blognames)
292         sys.stdout = stdout
293     drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
294     print "Done with dendrograms"
295     print "K=5"
296     kclust=kcluster(data, k=5)
297     print "K=10"
298     kclust=kcluster(data, k=10)
299     print "K=20"
300     kclust=kcluster(data, k=20)
301     coords=scaledown(data)
302     draw2d(coords, blognames, jpeg='blogs2d.jpg')

```

Listing 28: clusters.py

```
1 4827.50419137
2 3357.0495173
3 3337.77451337
4 3329.39524222
5 3325.08734145
6 3323.50421842
7 3323.33337939
8 3324.2545901
```

Listing 29: scaledown output

7 References

- [1] Internet Engineering Task Force (IETF). RFC-4287 The Atom Syndication Format. <https://tools.ietf.org/html/rfc4287>, 2005.
- [2] Toby Segaran. *Programming Collective Intelligence*. O'Reilly, first edition, 2007.