

Assignment 10

Fall 2014

CS595 Web Science

Dr. Michael Nelson

Mathew Chaney

December 11, 2014

Contents

1	Question 1	3
1.1	Question	3
1.2	Answer	3
2	Question 2	5
2.1	Question	5
2.2	Answer	5
3	Question 3	9
3.1	Question	9
3.2	Answer	9
4	Appendix A	11
5	References	15

Listings

1	matrix.py	4
2	docclass main	6
3	docclass main	9
4	docclass.py	11

List of Tables

1	Question 2: Predictions 1-50	7
2	Question 2: Predictions 51-100	8
3	Question 3: Assessments	10

1 Question 1

1.1 Question

Choose a blog or a newsfeed (or something similar as long as it has an Atom or RSS feed). It should be on a topic or topics of which you are qualified to provide classification training data. In other words, choose something that you enjoy and are knowledgeable of. Find a feed with at least 100 entries.

Create between four and eight different categories for the entries in the feed:

examples:

work, class, family, news, deals

liberal, conservative, moderate, libertarian

sports, local, financial, national, international, entertainment

metal, electronic, ambient, folk, hip-hop, pop

Download and process the pages of the feed as per the week 12 class slides.

1.2 Answer

To obtain the blog entries required for this assignment, the `matrix.py` script was again used to download the blog entries from “Kevin’s XL & Disc Golf Chronicles”, a blog written by Kevin Morrow about his motorcycling and disc golfing exploits. The categories I came up with for each blog entry are as follows:

1. game: recreational game(s) of disc golf
2. tourney: tournament round(s)
3. motorcycles: anything related to riding/owning motorcycles
4. event: community events/cookouts
5. diy: disc dyes/graphic design

The script downloaded entries from the atom feed [1] of the blog until a total of 100 entries were retrieved. It parsed each entry's title and saved them as a list to the `blog_content` file which will later be used for training the fisher classifier in Question 2.

```
1 import feedparser
2 import futures
3 import math
4 import md5
5 import re
6 import sys
7 import json
8
9 blog_uri = 'http://kevinmorrow.blogspot.com/feeds/posts/default'
10 data_file = 'blog_content'
11
12 def get_next(d):
13     for item in d.feed.links:
14         if item['rel'] == u'next':
15             return item['href']
16     return None
17
18 def parse_entries(entries, uri):
19     print('processing {}'.format(uri))
20     next = uri
21     while next is not None:
22         feed = feedparser.parse(next)
23         next = get_next(feed)
24         print('next {}'.format(next))
25         for entry in feed.entries:
26             if entry.title in entries:
27                 continue
28             entries.append(entry.title)
29             if len(entries) >= 100:
30                 next = None
31                 break
32     return entries
33
34 def load_data(filename):
35     entries = []
36     with open(filename) as infile:
37         return [entry.strip() for entry in infile]
38
39 if __name__ == '__main__':
40     old_entries = load_data(data_file)
41     entries = parse_entries(old_entries, blog_uri)
42     with open(data_file, 'w') as outfile:
43         for entry in entries:
44             outfile.write(entry + '\n')
```

Listing 1: matrix.py

2 Question 2

2.1 Question

Manually classify the first 50 entries, and then classify (using the fisher classifier) the remaining 50 entries. Report the `cprob()` values for the 50 titles as well. From the title or entry itself, specify the 1-, 2-, or 3-gram that you used for the string to classify. Do not repeat strings; you will have 50 unique strings. For example, in these titles the string used is marked with `*s`:

```
*Rachel Goswell* - "Waves Are Universal" (LP Review)
The *Naked and Famous* - "Passive Me, Aggressive You" (LP Review)
*Negativland* - "Live at Lewis's, Norfolk VA, November 21, 1992" (concert)
Negativland - "*U2*" (LP Review)
```

Note how "Negativland" is not repeated as a classification string.

Create a table with the title, the string used for classification, `cprob()`, predicted category, and actual category.

2.2 Answer

The `docclass.py` script was driven by the code shown in Listing 3. Before the script was run each of the 100 blog entries was classified manually to be used for training data later. This training is stored in the `training` file. The `docclass` main driver uses the first 50 entries as training for classifying the second 50 and then swaps each of these two sets. Tables 1 and 2 are the compiled training results with their actual classification and the `cprob` as calculated by the Fisher classifier. Refer to 4 in Appendix A for the full script.

```

206 entries = matrix.load_data(matrix.data_file)
207 cl = fisherclassifier(getwords)
208 cl.setdb('data.db')
209
210 T_HEAD = """\\begin{table}[h!]
211 \\centering
212 \\begin{tabular}{| l | l | l | l |}
213 \\hline
214 Entry Title & Actual & Predicted & cprob \\\\
215 \\hline
216 """
217
218 T_TAIL = """\\hline
219 \\end{tabular}
220 \\caption{Question 2: Predictions }
221 \\label{tab:mratings}
222 \\end{table}
223 """
224
225 def trainfrom(index=0):
226     keys = training.keys()
227     for key in keys[index:index+50]:
228         cl.train(key, training[key])
229     t = set(training.keys()[index:index+50])
230     k = set(entries)
231     rest = k - t
232     predict = {}
233     for item in rest:
234         group, prob = cl.classify(item)
235         predict[item] = (group, prob)
236     with open('predict' + str(index), 'w') as outfile:
237         outfile.write(T_HEAD)
238         for item, tup in predict.iteritems():
239             title = item.replace('&', '\\&').replace('#', '\\#')
240             row = ' & '.join([title, training[item], tup[0], str(tup[1])])
241             outfile.write(row + ' \\\\n')
242         outfile.write(T_TAIL)
243
244 if __name__ == '__main__':
245     with open('training') as infile:
246         training = {line.split('\\t')[0]: line.split('\\t')[1].strip() for line in infile}
247     trainfrom(0)
248     trainfrom(50)

```

Listing 2: docclass main

Entry Title	Actual	Predicted	cprob
Disc Girl	diy	diy	0.139554246962
2013 DGCR Hawk Hollow Weekend	event	tourney	0.135398562817
Hey! I'm Back!	personal	motorcycles	0.742810969879
Last Ride of the Season	motorcycles	motorcycles	0.569769009772
Promoting the Sport	event	news	0.215734958342
Dyeing For Some New Discs....	diy	personal	0.127742032263
It's Been a Weird Sportster Year	motorcycles	diy	0.459029096789
Frank Lloyd Wright Day Trip	news	game	0.519935002549
Mach III Re-fit....	diy	diy	0.327231205524
Mike Sale: The Quest for 2,500	event	news	0.348489346703
2014 Chili Cook-Off	tourney	motorcycles	0.568479771654
Betty Queen Open	tourney	tourney	0.28731988169
2013 Hawk Hollow Open: Ams	tourney	tourney	0.0528373296205
Feelin' Lucky? Punk!....	news	tourney	0.655185013039
Red Oak Rumble	tourney	tourney	0.655185013039
Beginners Guide for Disc Dyes (Long Post)	diy	diy	0.195308043476
Snow Round at Loriella	game	game	0.476013302099
PDGA World's, the rest of it...	tourney	tourney	0.230722681279
Great Way to Start the Year	tourney	news	0.466829424722
2013 Hawk Hollow Open... Pros	tourney	tourney	0.0802708824856
Latest Disc Dye...	diy	diy	0.0574532219591
Winter Round	game	tourney	0.59657359028
Sporty Surprise...	motorcycles	tourney	0.59657359028
Deluxe Retractable Birdie Bead Scoring System	diy	tourney	0.759831170994
Latest Dye...	diy	diy	0.154721589649
Seneca Sun Seeker	tourney	tourney	0.655185013039
2013 Loriella Challenge	tourney	tourney	0.257589575924
First Day in Charlottesville	tourney	game	0.476013302099
Bayville Bash IX	tourney	tourney	0.59657359028
State of the Sporty	motorcycles	news	0.327231205524
Mornin' Round at Loriella	game	game	0.476013302099
Hawk Hollow Open	tourney	tourney	0.0672478176833
2nd Annual LoCo Open	tourney	tourney	0.254631706407
SOMD Classic	tourney	tourney	0.59657359028
Winchester IFO	tourney	tourney	0.59657359028
D Day Ride...	motorcycles	game	0.38493019271
Virginia Team Invitational	event	tourney	0.166468597895
Good & Bad day of DG	game	game	0.257589575924
Vincent & Jules....	diy	tourney	0.59657359028
All Ready for the World's	news	diy	0.0772653501085
The Season is Over...	news	tourney	0.215734958342
A Full Day Saturday	game	game	0.257589575924
Last Day of Vacation	motorcycles	game	0.476013302099
Turkey Day Doubles	game	game	0.384502787693
New Improved Putter	diy	diy	0.215734958342
Fall is coming...	motorcycles	tourney	0.59657359028
Lost Another One...	diy	diy	0.327231205524
First rounds at the Worlds	tourney	news	0.351391490133
Ace Race Fun....	event	personal	0.476013302099
Day Two at the World's	tourney	diy	0.122142469344

Table 1: Question 2: Predictions 1-50

Entry Title	Actual	Predicted	cprob
Maryland vs Virginia Ice Bowl Battle IV	event	event	0.0979229226451
Had Some Fun This Morning...	personal	personal	0.0605701708213
2014 River City Open	tourney	tourney	0.108800373877
Saturday in Staunton	tourney	tourney	0.174085576264
Doin' a dye, dye, dye & dye...	diy	diy	0.0732870294308
Bored at work = evental disc dye	diy	diy	0.02046431997
Spotsy SuperDubs	tourney	tourney	0.23578679514
2014 Hawk Hollow Open	tourney	tourney	0.0204836924925
It's Been Awhile...	diy	diy	0.400936622169
Pretty Good Disc Golf Day	game	game	0.031850199739
Bored = Dye Some Plastic...	diy	diy	0.0290173583845
Almost There...	motorcycles	motorcycles	0.886142331508
All Hail the Disc!	diy	diy	0.0497016027828
Knocked for a Loop Today...	news	news	0.0717591100714
Orlando and Some Disc Golf	personal	personal	0.0190818350443
Disc Golf & Chili...	event	event	0.0257577881041
It's Been A While...	diy	diy	0.303191637893
Santa's Little Helper....	diy	diy	0.137680696132
Hawk Hollow Open - Ams	tourney	tourney	0.00985086638406
#830RatedforLife....	tourney	tourney	0.25
Blast From My Past	diy	diy	0.215734958342
Skyline Drive or Last Place?	motorcycles	motorcycles	0.981220963209
New DG Hobby... Para Cord	diy	diy	0.102330524752
Battle in the Blue Ridge	tourney	tourney	0.14611299113
My Friend is a World Champ	tourney	tourney	0.089580585114
Another Day... Another Dye	diy	diy	0.0629356658608
2013 Battlefield Open	tourney	tourney	0.0291686553403
'Merica, Fuck Yeah!	game	game	0.215734958342
Too Hot to Play...	game	game	0.166468597895
I won't be posting for a while...	personal	personal	0.0717591100714
Facebook is Making my World a Little Smaller...	diy	diy	0.0725730164948
Mach III Re-Paint...	diy	diy	0.127217607055
Some days are Just Better...	game	game	0.0979229226451
Multi-Color Disc Dyes	diy	diy	0.0252331502301
2013 Virginia Team Invitational	tourney	event	0.0345067265265
Hell Hath Frozen Over...	tourney	tourney	0.155662943557
My Best Buddy Died today...	personal	personal	0.155662943557
Gettin' Ready....	event	event	0.174085576264
2 Rounds @ Loriella	game	game	0.0849057427037
New Backpack	news	news	0.101483354394
Building a Course...	news	news	0.23578679514
Getting it Back Together...	motorcycles	motorcycles	0.81216172237
DGCR Mid-Atlantic Meet	event	event	0.155662943557
Lost a Friend This Week	personal	personal	0.0950728157762
Shaving Cream Disc Dyes	diy	diy	0.0440390153459
New Putter Dye...	diy	diy	0.0261190230258
Do a Little Dye, Play a Little Golf	event	event	0.0202847913416
First Ride of 2014	motorcycles	motorcycles	0.801415273277
Promoting the Club	news	news	0.155177975467
Multi-Color 2nd Attempt	diy	diy	0.0950728157762
Day Two at the World's	tourney	news	0.171383513075

Table 2: Question 2: Predictions 51-100

3 Question 3

3.1 Question

Assess the performance of your classifier in each of your categories by computing precision, recall, and F1. Note that the definitions of precisions and recall are slightly different in the context of classification

3.2 Answer

To calculate the *precision*, *recall* and *F-Measure* the `assess.py` script was used. This script parsed the pipe separated table stored in the file `predict_raw`, which contains all the predictions and cprob values for each item. The table is separated on each category.

```
1
2 def load_data(filename):
3     data = {}
4     with open(filename) as infile:
5         for line in infile:
6             entry, actual, predicted, cprob = line.split('|')
7             data[entry] = {'actual': actual, 'predicted': predicted, 'cprob': float(cprob.
8                             strip())}
9     return data
10
11 def assess(data, categories):
12     results = {}
13     for category in categories:
14         tp, fp, fn = float(0), float(0), float(0)
15         for entry, items in data.iteritems():
16             if data[entry]['actual'] != category:
17                 continue
18             if not data[entry]['predicted']:
19                 fn += 1
20             elif data[entry]['actual'] == data[entry]['predicted']:
21                 tp += 1
22             elif data[entry]['actual'] != data[entry]['predicted']:
23                 fp += 1
24         prec = tp / (tp + fp)
25         recall = tp / (tp + fn)
26         f1 = 2 * (prec * recall) / (prec + recall)
27         results[category] = {'p': str(prec), 'r': str(recall), 'f1': str(f1)}
28     return results
29
30 categories = ['game', 'tourney', 'motorcycles', 'event', 'diy']
31
32 T_HEAD = """\begin{table}[h!]
33 \centering
34 \begin{tabular}{|l|l|l|l|l|}
35 \hline
36 Category & Precision & Recall & F-Measure & \\
37 \hline
38 """
39
40 T_TAIL = """\hline
41 \end{tabular}
42 \caption{Question 3: Assessments }
43 \label{tab:assess}
44 \end{table}
45 """
46
47 data = load_data('predict_raw')
48 res = assess(data, categories)
49 with open('assess', 'w') as outfile:
50     outfile.write(T_HEAD)
51     for cat, table in res.iteritems():
52         outfile.write(' & '.join([cat, table['p'], table['r'], table['f1']]) + ' \\\n')
53     outfile.write(T_TAIL)
```

Listing 3: docclass main

Category	Precision	Recall	F-Measure
event	0.5	1.0	0.666666666667
game	0.909090909091	1.0	0.952380952381
tourney	0.785714285714	1.0	0.88
diy	0.884615384615	1.0	0.938775510204
motorcycles	0.454545454545	1.0	0.625

Table 3: Question 3: Assessments

4 Appendix A

```
1 from sqlite3 import dbapi2 as sqlite
2 import re
3 import math
4 import matrix
5
6 def getwords(doc):
7     splitter=re.compile('\W*')
8     # Split the words by non-alpha characters
9     words=[s.lower() for s in splitter.split(doc)
10            if len(s)>2 and len(s)<20]
11
12     # Return the unique set of words only
13     return dict([(w,1) for w in words])
14
15 class classifier:
16     def __init__(self, getfeatures, filename=None):
17         # Counts of feature/category combinations
18         self.fc={}
19         # Counts of documents in each category
20         self.cc={}
21         self.getfeatures=getfeatures
22
23     def setdb(self, dbfile):
24         self.con=sqlite.connect(dbfile)
25         self.con.execute('create table if not exists fc(feature,category,count)')
26         self.con.execute('create table if not exists cc(category,count)')
27
28     def incf(self, f, cat):
29         count=self.fc.count(f, cat)
30         if count==0:
31             self.con.execute("insert into fc values ('%s','%s',1)"
32                             % (f,cat))
33         else:
34             self.con.execute(
35                 "update fc set count=%d where feature='%s' and category='%s'"
36                 % (count+1,f,cat))
37
38     def fcount(self, f, cat):
39         res=self.con.execute(
40             'select count from fc where feature="%s" and category="%s"'
41             % (f,cat)).fetchone()
42         if res==None: return 0
43         else: return float(res[0])
44
45     def incc(self, cat):
46         count=self.catcount(cat)
47         if count==0:
48             self.con.execute("insert into cc values ('%s',1)" % (cat))
49         else:
50             self.con.execute("update cc set count=%d where category='%s'"
51                             % (count+1,cat))
52
53     def catcount(self, cat):
54         res=self.con.execute('select count from cc where category="%s"'
55                             % (cat)).fetchone()
56         if res==None: return 0
57         else: return float(res[0])
58
59     def categories(self):
60         cur=self.con.execute('select category from cc');
61         return [d[0] for d in cur]
62
63     def totalcount(self):
64         res=self.con.execute('select sum(count) from cc').fetchone();
65         if res==None: return 0
66         return res[0]
67
68     def train(self, item, cat):
69         features=self.getfeatures(item)
70         # Increment the count for every feature with this category
71         for f in features:
72             self.incf(f, cat)
```

```

75
76     # Increment the count for this category
77     self.incc(cat)
78     self.con.commit()
79
80     def fprob(self, f, cat):
81         if self.catcount(cat)==0: return 0
82
83         # The total number of times this feature appeared in this
84         # category divided by the total number of items in this category
85         return self.fcount(f, cat)/self.catcount(cat)
86
87     def weightedprob(self, f, cat, prf, weight=1.0, ap=0.5):
88         # Calculate current probability
89         basicprob=prf(f, cat)
90
91         # Count the number of times this feature has appeared in
92         # all categories
93         totals=sum([self.fcount(f, c) for c in self.categories()])
94
95         # Calculate the weighted average
96         bp=((weight*ap)+(totals*basicprob))/(weight+totals)
97         return bp
98
99
100
101
102     class naivebayes(classifier):
103
104         def __init__(self, getfeatures):
105             classifier.__init__(self, getfeatures)
106             self.thresholds={}
107
108         def docprob(self, item, cat):
109             features=self.getfeatures(item)
110
111             # Multiply the probabilities of all the features together
112             p=1
113             for f in features: p*=self.weightedprob(f, cat, self.fprob)
114             return p
115
116         def prob(self, item, cat):
117             catprob=self.catcount(cat)/self.totalcount()
118             docprob=self.docprob(item, cat)
119             return docprob*catprob
120
121         def setthreshold(self, cat, t):
122             self.thresholds[cat]=t
123
124         def getthreshold(self, cat):
125             if cat not in self.thresholds: return 1.0
126             return self.thresholds[cat]
127
128         def classify(self, item, default=None):
129             probs={}
130             # Find the category with the highest probability
131             max=0.0
132             for cat in self.categories():
133                 probs[cat]=self.prob(item, cat)
134                 if probs[cat]>max:
135                     max=probs[cat]
136                     best=cat
137
138             # Make sure the probability exceeds threshold*next best
139             for cat in probs:
140                 if cat==best: continue
141                 if probs[cat]*self.getthreshold(best)>probs[best]: return default
142             return best
143
144     class fisherclassifier(classifier):
145         def cprob(self, f, cat):
146             # The frequency of this feature in this category
147             clf=self.fprob(f, cat)
148             if clf==0: return 0
149
150             # The frequency of this feature in all the categories
151             freqsum=sum([self.fprob(f, c) for c in self.categories()])

```

```

152
153 # The probability is the frequency in this category divided by
154 # the overall frequency
155 p=clf/(freqsum)
156
157 return p
158 def fisherprob(self,item,cat):
159     # Multiply all the probabilities together
160     p=1
161     features=self.getfeatures(item)
162     for f in features:
163         p*=(self.weightedprob(f,cat,self.cprob))
164
165     # Take the natural log and multiply by -2
166     fscore=-2*math.log(p)
167
168     # Use the inverse chi2 function to get a probability
169     return self.invchi2(fscore,len(features)*2)
170 def invchi2(self,chi, df):
171     m = chi / 2.0
172     sum = term = math.exp(-m)
173     for i in range(1, df//2):
174         term *= m / i
175         sum += term
176     return min(sum, 1.0)
177 def __init__(self,getfeatures):
178     classifier.__init__(self,getfeatures)
179     self.minimums={}
180
181 def setminimum(self,cat,min):
182     self.minimums[cat]=min
183
184 def getminimum(self,cat):
185     if cat not in self.minimums: return 0
186     return self.minimums[cat]
187 def classify(self,item,default=None):
188     # Loop through looking for the best result
189     best=default
190     max=0.0
191     for c in self.categories():
192         p=self.fisherprob(item,c)
193         # Make sure it exceeds its minimum
194         if p>self.getminimum(c) and p>max:
195             best=c
196             max=p
197     return best, p
198
199 def sampletrain(cl):
200     cl.train('Nobody owns the water.','good')
201     cl.train('the quick rabbit jumps fences','good')
202     cl.train('buy pharmaceuticals now','bad')
203     cl.train('make quick money at the online casino','bad')
204     cl.train('the quick brown fox jumps','good')
205
206 entries = matrix.load_data(matrix.data_file)
207 cl = fisherclassifier(getwords)
208 cl.setdb('data.db')
209
210 T_HEAD = """\\begin{table}[h!]
211 \\centering
212 \\begin{tabular}{|l|l|l|l|l|}
213 \\hline
214 Entry Title & Actual & Predicted & cprob & \\\\
215 \\hline
216 """
217
218 T_TAIL = """\\hline
219 \\end{tabular}
220 \\caption{Question 2: Predictions }
221 \\label{tab:mratings}
222 \\end{table}
223 """
224
225 def trainfrom(index=0):
226     keys = training.keys()
227     for key in keys[index:index+50]:
228         cl.train(key, training[key])

```

```

229 t = set(training.keys()[index:index+50])
230 k = set(entries)
231 rest = k - t
232 predict = {}
233 for item in rest:
234     group, prob = cl.classify(item)
235     predict[item] = (group, prob)
236 with open('predict' + str(index), 'w') as outfile:
237     outfile.write(T_HEAD)
238     for item, tup in predict.iteritems():
239         title = item.replace('&', '\\&').replace('#', '\\#')
240         row = ' & '.join([title, training[item], tup[0], str(tup[1])])
241         outfile.write(row + ' \\\n')
242     outfile.write(T_TAIL)
243
244 if __name__ == '__main__':
245     with open('training') as infile:
246         training = {line.split('\t')[0]: line.split('\t')[1].strip() for line in infile}
247     trainfrom(0)
248     trainfrom(50)

```

Listing 4: docclass.py

5 References

- [1] Internet Engineering Task Force (IETF). RFC-4287 The Atom Syndication Format. <https://tools.ietf.org/html/rfc4287>, 2005.